

---

## Project 3 Guidelines: Implementation of Neural Networks for Object Detection

This project is structured to integrate neural networks in object detection. The objective is to facilitate practical comprehension and application of neural network concepts, from classification to object detection tasks.

### Part 1: Report on Classification Model Backbone (10 Points)

#### Backbone Architecture Selection

- Aggregate the IDs of all group members, and utilize the last digit to select your backbone architecture:
  - 0-3: ResNet18
  - 4-6: VGG16
  - 7-9: MobileNet V3

#### Requirements

1. Document the process of summing the IDs and selecting the architecture. A 10 point deduction will apply for incorrect ID sum calculations. Example:  $id1=12345678$ ,  $id2=87654321 \rightarrow 1+2+3+\dots+8+7+6+\dots=72 \rightarrow 7+2=9 \rightarrow$  MobileNet V3.
2. Present inference results from the pre-trained network to demonstrate its current capabilities on a few random images.
3. Conduct a detailed analysis (minimum two pages) on the chosen architecture and inference results, focusing on unfamiliar blocks or components not covered in course lectures. Utilize original research papers and additional resources for this analysis.

### Part 2: Single Class, Single Object Per Frame- Basic Object Detection Model Implementation (60 Points)

In this section, you are tasked with developing an object detection system, leveraging a preselected classification model as the foundational backbone. This endeavor takes us beyond the confines of our course syllabus, venturing into the application of transfer learning from established classification frameworks to the nuanced domain of object detection. A crucial aspect of this project involves your engagement with independent research to comprehend and integrate each essential component of an object detection model.

---

## What is a backbone?

Before delving into the specific tasks, it's imperative to understand the significance of utilizing a pretrained classification model as the backbone for your object detection system. Pretrained classification models, such as ResNet, VGG, or MobileNet, have been extensively trained on vast datasets like ImageNet, enabling them to develop a rich, hierarchical representation of visual features. This foundational knowledge facilitates the model's ability to discern complex patterns and features in new images, a capability that is immensely beneficial for the task of object detection. By employing these pretrained networks, you capitalize on their learned representations, effectively shortcutting the need for training a model from scratch. This approach not only saves considerable computational resources and time but also enhances the model's detection capabilities, especially when dealing with limited data for the specific task at hand. The transfer of learned features from classification to detection tasks is a testament to the versatility and efficiency of deep learning models, allowing for sophisticated applications such as accurate and robust object detection.

## Model Construction and Evaluation

1. **Architecture Development:** Construct an object detection architecture utilizing the selected classification backbone. This phase includes several critical steps, each necessitating prior investigation and detailed justification:
  - **Model Building:** Using only the pretrained backbone + basic torch building blocks (e.g: conv2d, relu, etc...), build a model
  - **Axis-Aligned Bounding Boxes (AABB):** Investigate AABB usage in object detection for a single label type.
  - **Loss Functions and Metrics:** Explore various loss functions and metrics relevant to object detection. Select the option that best suit the specifics of your dataset and model objectives, justifying your choices based on the theoretical advantages they provide for enhancing model accuracy and reliability.
  - **Data Partitioning:** Examine strategies for splitting your dataset into training, validation, and testing sets. Discuss how this strategy will influence model generalization and performance evaluation.
  - **Data Augmentation:** Investigate and implement data augmentation techniques suitable for object detection, using [albumentations](#) or [torchvision](#). Detail the augmentations chosen and discuss their expected impact on model robustness and generalization.
  - **TensorBoard Monitoring:** Use [tensorboard](#) to monitor training and validation progress, focusing on metrics relevant AABB accuracy. Discuss the importance of these visualizations in

---

optimizing your model.

2. **Dataset Selection:** Select an appropriate dataset for object detection (a good place for it can be Roboflow Public Datasets). Justify your selection based on the dataset's relevance, complexity, and ability to challenge and validate your model's effectiveness.
3. **Model Training:** After constructing your model and selecting a dataset, proceed to train your object detection system. This stage is crucial for fine-tuning the model's parameters to effectively recognize and localize the objects of interest within the dataset. Provide a detailed account of your training process, including:
  - Configuration of training parameters (e.g., learning rate, batch size).
  - Strategies employed to avoid overfitting and ensure generalization.
  - Analysis of training and validation loss curves to monitor the learning process.
  - Adjustments made based on the performance metrics observed during training.
4. **Inference on External Video:** Test your model on a video not included in the training set to evaluate its real-world applicability. Document your methodology, outcomes, and any challenges encountered, offering insights into the model's operational performance. **NOTE: This step is mandatory. Failing to send a final video, different than the train dataset, with good results will impact heavily on your final grade of this part.** Examples of points deduction: Video is too "easy" (white background, static object, minimal changes); Partial detection of some frames (small objects, rotated objects).

Each section of your report should meticulously document your research, decision-making process, and implementation approach. This comprehensive methodology will not only showcase your technical expertise but also your capacity to critically engage with novel information and apply it to practical scenarios.

### Part 3: Single Class, Multiple Objects Per Frame (10 Points)

In this part, you will extend your object detection model to handle scenarios where multiple objects of the same class are present in a single frame. This step is essential to demonstrate your model's ability to detect and localize multiple instances accurately.

#### Key Objectives:

1. **Data Preparation and Augmentation:**
  - Modify your dataset to include images with multiple objects of the same class in each frame.

- 
- Apply suitable data augmentation techniques (e.g., scaling, flipping, rotation, and cropping) to enhance model robustness and generalization.

## 2. **Model Adjustment:**

- Update your object detection model to handle multiple bounding boxes for the same class within a single image.
- Ensure the model can output multiple predictions, including bounding box coordinates and confidence scores, for each detected object.

## 3. **Loss Function Adaptation:**

- Adapt your loss function to consider multiple objects in the same frame.
- Implement Non-Maximum Suppression (NMS) to handle overlapping predictions effectively and avoid duplicate detections.

## 4. **Evaluation Metrics:**

- Use metrics like Intersection over Union (IoU) and Mean Average Precision (mAP) to evaluate the model's ability to detect and localize multiple objects accurately.

## 5. **Training and Validation:**

- Train your model on the modified dataset.
- Validate the model's performance using a separate validation set with multiple objects per frame.
- Monitor training and validation metrics using [TensorBoard](#) or a similar tool.

## 6. **Inference on External Video:**

- Test your model on a new video containing multiple objects of the same class in each frame.
- Document the inference process and results, showcasing the model's capability to handle multiple objects.

# **Part 4: Multiple Classes, Multiple Objects Per Frame (10 Points)**

In this section, you are tasked with advancing your object detection model to handle scenarios involving multiple classes and multiple objects within a single frame. This step is critical for developing a robust detection system capable of recognizing and localizing various objects across diverse environments. You will integrate and refine your model, leveraging the classification backbone chosen in previous sections, to accurately detect and classify multiple objects in complex scenes.

Key objectives include:

- 
1. **Data Preparation and Augmentation:** Change your dataset to include images with multiple objects and classes. Apply data augmentation techniques to improve model generalization.
  2. **Model Adjustment:** Modify your model architecture to support multi-object, multi-class detection. This may involve integrating additional layers or adjusting the existing structure to process and predict multiple objects simultaneously.
  3. **Loss Function Adaptation:** Implement or modify loss functions that can effectively handle the complexity of multiple classes and objects, such as Cross-Entropy Loss for classification and a combination of losses for localization.
  4. **Evaluation Metrics:** Utilize appropriate metrics for multi-object detection, including Mean Average Precision (mAP) for performance evaluation across various classes and Intersection over Union (IoU) for assessing localization accuracy.
  5. **Training and Validation:** Conduct thorough training sessions, using a split dataset for validation to monitor and enhance model performance iteratively.

**NOTE: You will lose points if you fail to show the model capabilities to work on different object classes in the same frame**

## **Part 5: Oriented Bounding Box Implementation (10 Points)**

This section introduces the implementation of Oriented Bounding Boxes (OBBs) to your object detection model, using the selected classification backbone. OBBs provide a more accurate localization by allowing bounding boxes to rotate, closely fitting the object's actual orientation and shape. This advanced feature is crucial for objects with varied poses and angles, enhancing detection precision beyond the capabilities of Axis-Aligned Bounding Boxes (AABBs).

**NOTE: You will lose points if you fail to show the model capabilities to work on different object classes in the same frame**

### **Implementation Overview**

- **Study OBB Concepts:** Briefly research the fundamentals of OBBs, focusing on their geometric properties and how they offer a more precise encapsulation of objects by including orientation.
- **Model Adaptation:** Similar to Part 2, modify your detection model to predict OBB parameters: center coordinates, dimensions, and rotation angle. This involves adjusting the architecture to output these additional details, akin to how you structured your model for AABB predictions.

- 
- **Loss Function and Metrics:** Adapt your loss functions and evaluation metrics to accommodate the complexities of OBBs, especially for accurately predicting and evaluating the orientation of objects.
  - **Training and Evaluation:** Train your model on an appropriate dataset, ensuring it includes orientation annotations. The training process and evaluation should mirror the approach taken in Part 2, with adjustments made for the unique aspects of OBB prediction.

This concise task aims to enhance your model's detection capabilities by incorporating orientation, a critical step for achieving higher precision in object detection.

### Submission Guidelines

- Form groups of up to two members.
- Compile results in a `.zip` file named `PROJ3_NAME1_ID1_NAME2_ID2.zip`, containing:
  - A comprehensive summary of the project, including methodologies, assumptions, successes, and limitations of the algorithm.
  - Source code in `.py` format.
  - Output videos in a supported format.
  - Include the final result video either as a file within the `.zip` or provide a YouTube link in the report (ensure visibility at both the beginning and end of the report).

### Deadline

- Submit by the day preceding the start of the next semester.

Good Luck! Yoni Chechik